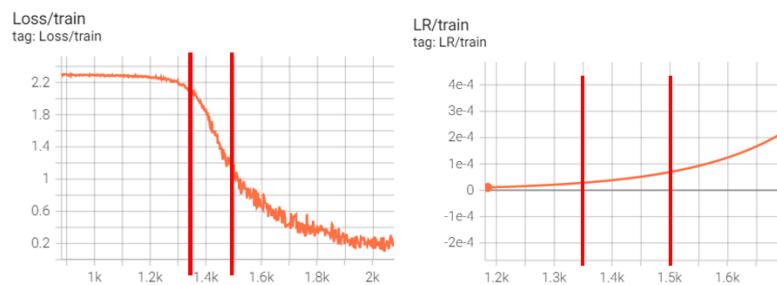


## CS 5260 Assignment 6

Jianxin Bi, A0219018J

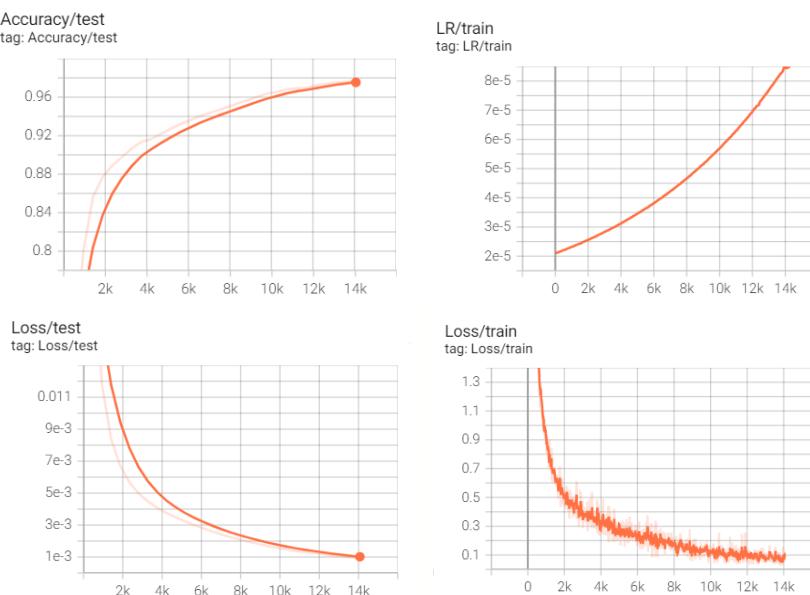
[https://github.com/bijianxin292430887/cs5260\\_assignment6](https://github.com/bijianxin292430887/cs5260_assignment6)

1. Optimizer: Adam( $lr=0.001$ ,  $\text{betas} = (0.9, 0.999)$ ,  $\text{eps}=1e-08$ ,  $\text{weight\_decay}=0$ )
2. Lr\_scheduler:
  - a. LambdaLR
  - b. ExponentialLR
3. LR range test for Adam optimizer: lr should take within  $[2.1e-5, 6.8e-5]$ , where the loss decrease the fastest. Test plots with LR region highlighted with red lines are shown below:



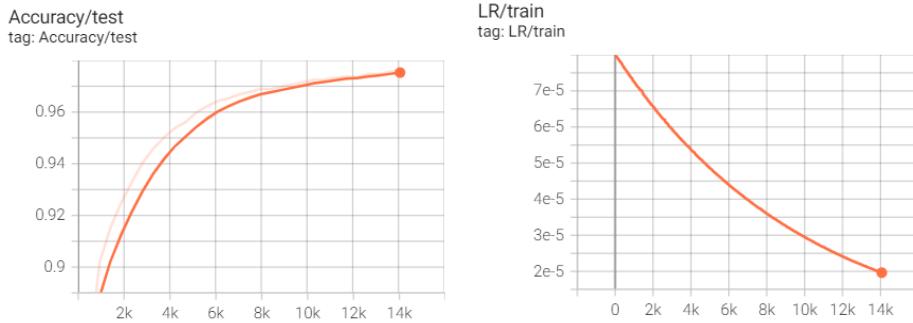
Training with Adam + ExponentialLR:

1. `lr_scheduler = torch.optim.lr_scheduler.ExponentialLR(optimizer, gamma=1.0001)`
2. `optimizer = torch.optim.Adam(model.parameters(), lr=2.1e-5, betas=(0.9, 0.999), eps=1e-08, weight_decay=0)`
3. Achieve **97.7%** accuracy with default LeNet5 hyperparameters.  
Batchsize = 128, Epochs = 30.



If we make the learning rate decrease along training, where  
1. `lr_scheduler = torch.optim.lr_scheduler.ExponentialLR(optimizer, gamma=0.9999)`

2. `optimizer = torch.optim.Adam(model.parameters(), lr=8e-5, betas=(0.9, 0.999), eps=1e-08, weight_decay=0)`
3. We obtain test accuracy **97.6%** surprisingly.



### Training with Adam + OneCycleLR:

1. `optimizer = torch.optim.Adam(model.parameters(), lr=8e-5, betas=(0.9, 0.999), eps=1e-08, weight_decay=0)`
2. `lr_scheduler = torch.optim.lr_scheduler.OneCycleLR(optimizer, max_lr=8e-5, steps_per_epoch = len(train_dataloader), epochs = 30)`
3. Achieve **97.26%** accuracy with default LeNet5 hyperparameters.  
Batchsize = 128, Epochs = 30.

